


Earlier custom Python model · full four-point sweep

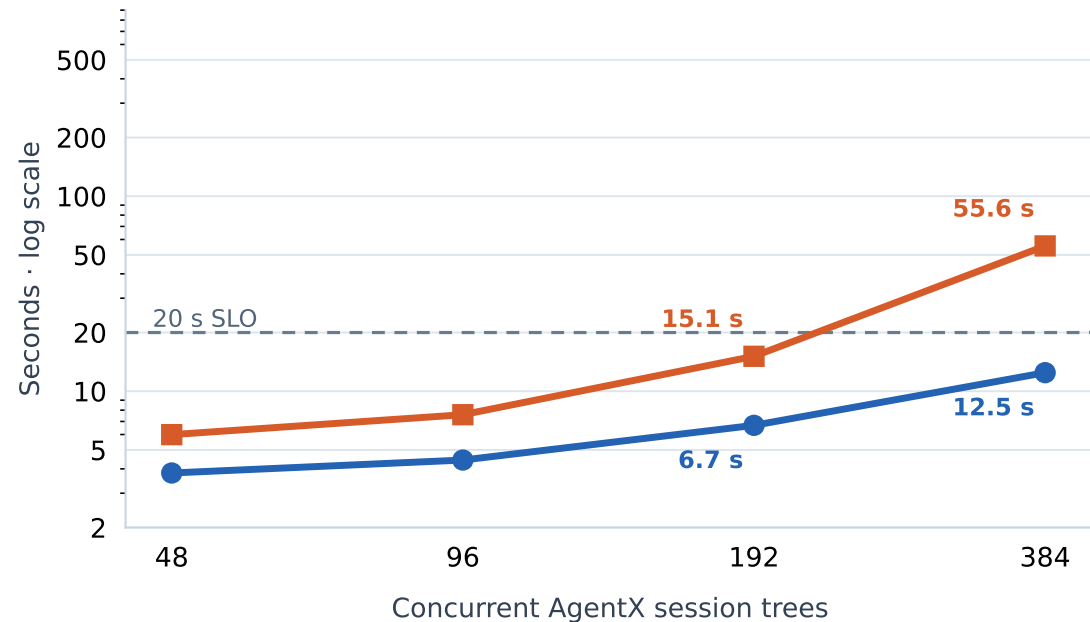
Actual AIPerf workload replay · custom fixed-rate serving · 28.4M-token cache per worker

—●— KV —■— RR

Throughput



P95 time to first token



Both modeled throughput curves are still rising at 384: no knee is established within this range.

This custom model is not NVIDIA DynoSim; its missing saturation and optimistic latency limit hardware/SLO decisions.