


# Native DynoSim V10 · throughput and latency

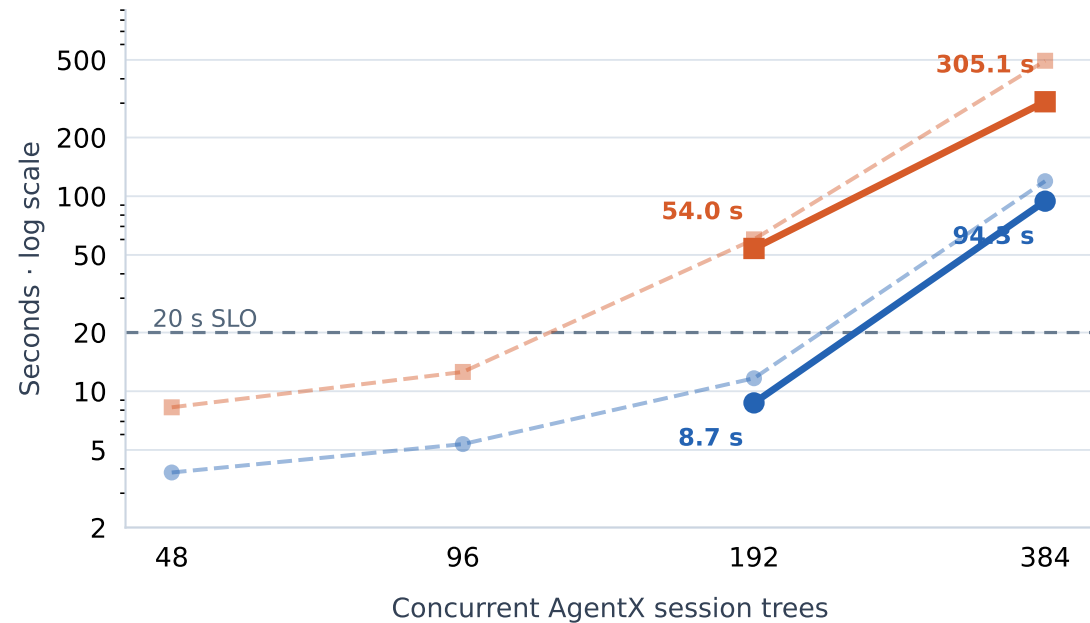
Dynamo 1.4.2 + local V10 patches · AIC 0.11.0 · 24-GPU model · normal AIPerf HTTP replay

—●— KV hardware    —■— RR hardware    —●— KV V10    —■— RR V10

## Throughput



## P95 time to first token



Solid: V10 at 192 and 384 only. Faint dashed: hardware. Native 48/96 points are unavailable.

Throughput passes the recorded  $\pm 20\%$  gate; TTFT remains uncalibrated. The exact knee needs more samples.