


Native DynoSim V10 · throughput and latency


Dynamo 1.4.2 + local V10 patches · AIC 0.11.0 · 24-GPU model · normal AIPerf HTTP replay

—●— KV hardware —●— KV V10 —◆— KV scale 3 / credit 0.8 —✱— KV temperature 0.5
—□— RR hardware —□— RR V10 —▲— KV scale 2 / credit 0.8

Throughput



P95 time to first token



KV 192→384: throughput -7.4%, TTFT p95 ×10.8. RR 96→192: throughput +11.6%, TTFT p95 ×4.4.

Solid: V10; faint dashed: hardware. ★: sampled default-policy throughput peaks; not SLO choices.

Four calibration points + eight holdouts, one frozen build. Tuned knees remain unresolved; latency requires validation.