



Evidence behind the latency rise · measured hardware

Prefix reuse declines with offered population



Live sessions and active requests are different



Cache ratio comes from frontend histogram sums, not per-engine cache scrapes. In-flight includes failed requests and is time-weighted over each profiling request envelope. These signals support a prefill/queue hypothesis, not a complete bottleneck proof.