



Measured KV-router flag sweep · concurrency 480 · same 64-GPU fleet


Total input + output tokens/s/GPU




TTFT p95 (seconds; lower is better)



E2E I90 (output tok/s/user; higher is better)



Frontend-reported cached input (%)



Credit 1.5 vs default: TTFT p95 −15.4%, I90 +13.1%, throughput +0.3%; repeats are needed.
All four KV runs have zero exported errors. RR has 39 client errors / 8,317 profiling records and is an overload reference.