


Agg · 24 GPUs · selected KV and RR throughput

Highest measured throughput per policy passing TTFT p95 <10 s and E2E I90 ≥20 output tok/s/user



C = live sessions. Load scale = prefill load scale; credit = KV overlap credit. KV uses temperature 0, decay 0 and FCFS. Throughput includes cached input and output tokens. Ratios use the selected RR run; concurrency differs by policy. Cache hit = client-reported cached-input percentage. All selected runs pass the queue check and have zero client errors.