



Matched disagg checks · 12P+6D / 72 GPUs · KV routing only

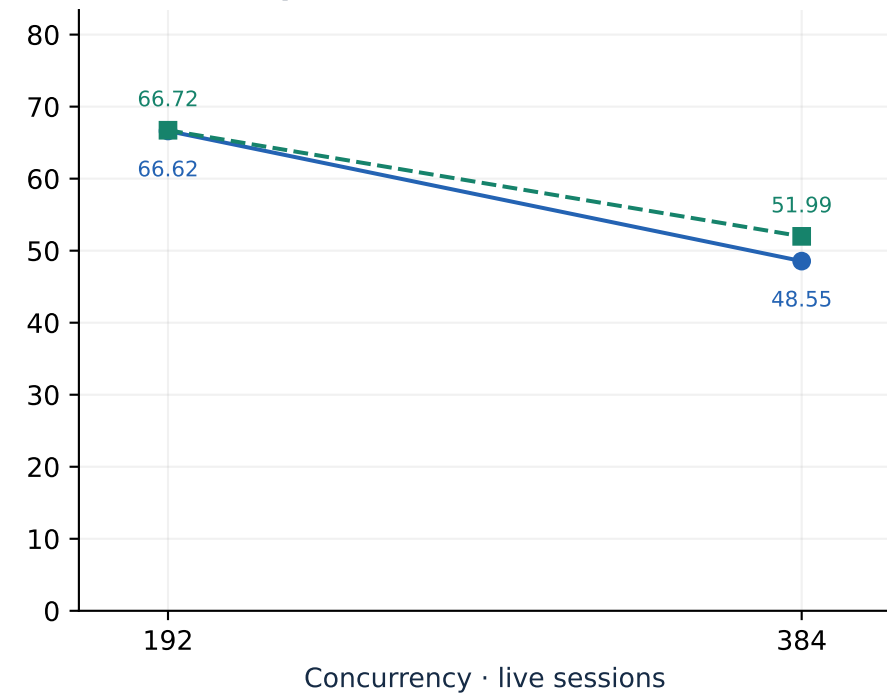
Total input + output tok/s/GPU



TTFT p95 (seconds)



E2E I90 (output tok/s/user)



Same topology and concurrency, matched warmup inputs, one-hour profiles; native V11 with frozen V10 timing.
Two historical reference points, not the current 64-GPU D88 fleet or an RR validation. Lines connect samples; they do not locate a knee.